

# Pre-Announcements

---

Your friend from this class is here to tell you about an event.

- Sutardja center for Entrepreneurship is holding a hackathon called Cal Innovates.
- April 14th/15th [?]. Open to anybody, CS, econ, business, premed, etc.
- Recruiters from all sorts of companies will be there.
- Prizes will be there, not just electronics gadgets, but also ability to pitch to Skydeck which is a venture capital firm in house at Berkeley.
- Register by this Thursday for raffle eligibility.
- Innovateatcal.com

# Announcements

---

Proj3: It's out. If you haven't started yet, you are behind.

# CS61B

## Lecture 32: Sorting I

- The Sorting Problem
- Selection Sort, Heapsort
- Mergesort
- Insertion Sort
- Shell's Sort (Extra)



# Sorting - Definitions

---

A sort is a permutation (re-arrangement) of a sequence of elements that brings them into order according to some **total order**. A total order  $\leq$  is:

- Total:  $x \leq y$  or  $y \leq x$  for all  $x, y$
- Reflexive:  $x \leq x$
- Antisymmetric:  $x \leq y$  and  $y \leq x$  iff  $x = y$  ( $x$  and  $y$  are equivalent).
- Transitive:  $x \leq y$  and  $y \leq z$  implies  $x \leq z$ .

In Java, total order is typically specified by `compareTo` or `compare` methods.

- May be inconsistent with `equals`! For example sorting an array of `Strings` by length has items that are equivalent, but not equal, e.g. “cat” and “dog”.

# Sorting: An Alternate Viewpoint

---

An ***inversion*** is a pair of elements that are out of order.



Yoda

0 1 1 2 3 4 8 6 9 5 7

8-6 8-5 8-7 6-5 9-5 9-7

(6 inversions out of 55 max)



Gabriel Cramer

Goal of sorting:

- Given a sequence of elements with Z inversions.
- Perform a sequence of operations that reduces inversions to 0.

# Performance Definitions

---

Characterizations of the runtime efficiency are sometimes called the **time complexity** of an algorithm. Examples:

- DFS has time complexity  $\Theta(V+E)$ .

Characterizations of the “extra” memory usage of an algorithm is sometimes called the **space complexity** of an algorithm.

- DFS has space complexity  $\Theta(V)$ .
  - Note that the graph takes up space  $\Theta(V+E)$ , but we don't count this as part of the runtime of DFS, since we're only accounting for the extra space that DFS uses.

# Selection Sort and Heapsort

# Selection Sort

---

We've seen this already.

- Find smallest item.
- Swap this item to the front and 'fix' it.
- Repeat for unfixed items until all items are fixed.
- Demo: <https://goo.gl/g14Cit>

Sort Properties:

- $\Theta(N^2)$  time if we use an array (or similar data structure).

Seems inefficient: We look through entire remaining array every time to find the minimum.



# Naive Heapsort: Leveraging a Max-Oriented Heap

---

Idea: Instead of rescanning entire array looking for minimum, maintain a heap so that getting the minimum is fast!

For reasons that will become clear soon, we'll use a max-oriented heap.

A min heap would work as well, but wouldn't be able to take advantage of the fancy trick in a few slides.

Naive heapsorting N items:

- Insert all items into a max heap, and discard input array. Create output array.
- Repeat N times:
  - Delete largest item from the max heap.
  - Put largest item at the end of the unused part of the output array.

Naive Heapsort Demo: <https://goo.gl/EZWwSJ>

## Naive Heapsort Runtime: <http://yellkey.com/suffer>

---

Naive heapsorting N items:

- Insert all items into a max heap, and discard input array. Create output array.
- Repeat N times:
  - Delete largest item from the max heap.
  - Put largest item at the end of the unused part of the output array.

What is the TOTAL runtime of naive heapsort?

- A.  $\Theta(N)$
- B.  $\Theta(N \log N)$
- C.  $\Theta(N^2)$ , but faster than selection sort.

# Heapsort Runtime Analysis

---

Use the magic of the heap to sort our data.

- Getting items into the heap  $O(N \log N)$  time.
- Selecting *largest* item:  $\Theta(1)$  time.
- Removing *largest* item:  $O(\log N)$  for each removal.

Informally?? [good eye!]

Informally, overall runtime is  $O(N \log N) + \Theta(N) + O(N \log N) = \mathbf{O(N \log N)}$

- Far better than selection sort!

Memory usage is  $\Theta(N)$  to build the additional copy of all of our data.


- Worse than selection sort, but probably no big deal (??).
- Can eliminate this extra memory cost with same fancy trickery.

# In-place Heapsort

---

Alternate approach, treat input array as a heap!

- Rather than inserting into a new array of length  $N + 1$ , use a process known as “bottom-up heapification” to convert the array into a heap.
  - To bottom-up heapify, just sink nodes in reverse level order.
- Avoids need for extra copy of all data.
- Once heapified, algorithm is almost the same as naive heap sort.



For this algorithm we don't leave spot 0 blank.

In-place heap sort: [Demo](#)

## In-place Heapsort Runtime: <http://yellkey.com/fail>

---

Use the magic of the heap to sort our data.

- Bottom-up Heapification:  $O(???)$  time.
- Selecting *largest* item:  $\Theta(1)$  time.
- Removing *largest* item:  $O(\log N)$  for each removal.

Give the time complexity of in-place heapsort in big O notation.

- A.  $O(N)$
- B.  $O(N \log N)$
- C.  $O(N^2)$

# In-place Heapsort Runtime

---

Use the magic of the heap to sort our data.

- Bottom-up Heapification:  $O(N \log N)$  time.
- Selecting *largest* item:  $\Theta(1)$  time.
- Removing *largest* item:  $O(\log N)$  for each removal.

Give the time complexity of in-place heapsort in big O notation.

**A.  $O(N \log N)$**

Bottom-up heapification is  $N$  sink operations, each taking no more than  $O(\log N)$  time, so overall runtime for heapification is  $O(N \log N)$ .

- Extra for experts, show that bottom-up heapification is  $\Theta(N)$  in the worst case.
- More extra for experts, show heapsort is  $\Theta(N \log N)$  in the worst case.

## In-place Heapsort: <http://yellkey.com/once>

---

What is the **memory complexity** of Heapsort?

- Also called “space complexity”.
- A.  $\Theta(1)$
- B.  $\Theta(\log N)$
- C.  $\Theta(N)$
- D.  $\Theta(N \log N)$
- E.  $\Theta(N^2)$

# In-place Heapsort

---

What is the **memory complexity** of Heapsort?

- A.  $\Theta(1)$
- B.  $\Theta(\log N)$
- C.  $\Theta(N)$
- D.  $\Theta(N \log N)$
- E.  $\Theta(N^2)$

The only extra memory we need is a constant number instance variables, e.g. size.

- Unimportant caveat: If we employ recursion to implement various heap operations, space complexity is  $\Theta(\log N)$  due to need to track recursive calls. The difference between  $\Theta(\log N)$  and  $\Theta(1)$  space is effectively nothing.



# Sorts So Far

---

	Best Case Runtime	Worst Case Runtime	Space	Demo	Notes
<a href="#">Selection Sort</a>	$\Theta(N^2)$	$\Theta(N^2)$	$\Theta(1)$	<a href="#">Link</a>	
<a href="#">Heapsort</a> (in place)	$\Theta(N)^*$	$\Theta(N \log N)$	$\Theta(1)$	<a href="#">Link</a>	Bad cache (61C) performance.

\*: An array of all duplicates yields linear runtime for heapsort.

# Mergesort

# Mergesort

---

We've seen this one before as well.

Mergesort: [Demo](#)

- Split items into 2 roughly even pieces.
- Mergesort each half (steps not shown, this is a recursive algorithm!)
- Merge the two sorted halves to form the final result.

Time complexity, [analysis from previous lecture](#):  $\Theta(N \log N)$  runtime)

- Space complexity with aux array: Costs  $\Theta(N)$  memory.

Also possible to do in-place merge sort, but algorithm is very complicated, and runtime performance suffers by a significant constant factor.

# Sorts So Far

---

	Best Case Runtime	Worst Case Runtime	Space	Demo	Notes
<a href="#">Selection Sort</a>	$\Theta(N^2)$	$\Theta(N^2)$	$\Theta(1)$	<a href="#">Link</a>	
<a href="#">Heapsort</a> (in place)	$\Theta(N)^*$	$\Theta(N \log N)$	$\Theta(1)$	<a href="#">Link</a>	Bad cache (61C) performance.
<a href="#">Mergesort</a>	$\Theta(N \log N)$	$\Theta(N \log N)$	$\Theta(N)$	<a href="#">Link</a>	Faster than heap sort.

\*: An array of all duplicates yields linear runtime for heapsort.

# Insertion Sort

# Insertion Sort

---

General strategy:

- Starting with an empty output sequence.
- Add each item from input, inserting into output at right point.

Naive approach, build entirely new output: [Demo \(Link\)](#)

Input:

32	15	2	17	19	26	41	17	17
----	----	---	----	----	----	----	----	----

Output:

# Insertion Sort


---

General strategy:

- Starting with an empty output sequence.
- Add each item from input, inserting into output at right point.

Naive approach, build entirely new output: [Demo \(Link\)](#)

Input:



Output:



# Insertion Sort

---

General strategy:

- Starting with an empty output sequence.
- Add each item from input, inserting into output at right point.

For naive approach, if output sequence contains  $k$  items, worst cost to insert a single item is  $k$ .

- Might need to move everything over.

More efficient method:

- Do everything in place using swapping: [Demo \(Link\)](#)



# In-place Insertion Sort

Two more examples.

7 swaps:

P	O	T	A	T	O	
P	O	T	A	T	O	(0 swaps)
O	P	T	A	T	O	(1 swap)
O	P	T	A	T	O	(0 swaps)
A	O	P	T	T	O	(3 swaps)
A	O	P	T	T	O	(0 swaps)
A	O	P	T	T	O	(3 swaps)

**Purple:** Element that we're moving left (with swaps).

**Black:** Elements that got swapped with purple.

**Grey:** Not considered this iteration.

36 swaps:

S	O	R	T	E	X	A	M	P	L	E	
S	O	R	T	E	X	A	M	P	L	E	(0 swaps)
O	S	R	T	E	X	A	M	P	L	E	(1 swap)
O	R	S	T	E	X	A	M	P	L	E	(1 swap)
O	R	S	T	E	X	A	M	P	L	E	(0 swaps)
E	O	R	S	T	X	A	M	P	L	E	(4 swaps)
E	O	R	S	T	X	A	M	P	L	E	(0 swaps)
A	E	O	R	S	T	X	M	P	L	E	(6 swaps)
A	E	M	O	R	S	T	X	P	L	E	(5 swaps)
A	E	M	O	P	R	S	T	X	L	E	(4 swaps)
A	E	L	M	O	P	R	S	T	X	E	(7 swaps)
A	E	L	M	O	P	R	S	T	X	E	(8 swaps)

# Insertion Sort Runtime: <http://yellkey.com/sing>

What is the runtime of insertion sort?

- A.  $\Omega(1)$ ,  $O(N)$
- B.  $\Omega(N)$ ,  $O(N)$
- C.  $\Omega(1)$ ,  $O(N^2)$
- D.  $\Omega(N)$ ,  $O(N^2)$
- E.  $\Omega(N^2)$ ,  $O(N^2)$

36 swaps:

S	O	R	T	E	X	A	M	P	L	E		
S	O	R	T	E	X	A	M	P	L	E	(0 swaps)	
O	S	R	T	E	X	A	M	P	L	E	(1 swap)	
O	R	S	T	E	X	A	M	P	L	E	(1 swap)	
O	R	S	T	E	X	A	M	P	L	E	(0 swaps)	
E	O	R	S	T	E	X	A	M	P	L	E	(4 swaps)
E	O	R	S	T	X	A	M	P	L	E	(0 swaps)	
A	E	O	R	S	T	X	M	P	L	E	(6 swaps)	
A	E	M	O	R	S	T	X	P	L	E	(5 swaps)	
A	E	M	O	P	R	S	T	X	L	E	(4 swaps)	
A	E	L	M	O	P	R	S	T	X	E	(7 swaps)	
A	E	L	M	O	P	R	S	T	X		(8 swaps)	

# Insertion Sort Runtime: <http://yellkey.com/sing>

What is the runtime of insertion sort?

- A.  $\Omega(1)$ ,  $O(N)$
- B.  $\Omega(N)$ ,  $O(N)$
- C.  $\Omega(1)$ ,  $O(N^2)$
- D.  **$\Omega(N)$ ,  $O(N^2)$**
- E.  $\Omega(N^2)$ ,  $O(N^2)$

You may recall  $\Omega$  is not “best case”.

So technnnnniically you could also say  
 $\Omega(1)$

36 swaps:

S	O	R	T	E	X	A	M	P	L	E		
S	O	R	T	E	X	A	M	P	L	E	(0 swaps)	
O	S	R	T	E	X	A	M	P	L	E	(1 swap )	
O	R	S	T	E	X	A	M	P	L	E	(1 swap )	
O	R	S	T	E	X	A	M	P	L	E	(0 swaps)	
E	O	R	S	T	E	X	A	M	P	L	E	(4 swaps)
E	O	R	S	T	E	X	A	M	P	L	E	(0 swaps)
A	E	O	R	S	T	E	X	M	P	L	E	(6 swaps)
A	E	M	O	R	S	T	E	X	P	L	E	(5 swaps)
A	E	M	O	P	R	S	T	E	X	L	E	(4 swaps)
A	E	L	M	O	P	R	S	T	E	X		(7 swaps)
A	E	L	M	O	P	R	S	T	E	X		(8 swaps)

## Picking the Best Sort: <http://yellkey.com/national>

---

Suppose we do the following:

- Read 1,000,000 integers from a file into an array of length 1,000,000.
- Mergesort these integers.
- Select one integer randomly and change it.
- Sort using algorithm X of your choice.

Which sorting algorithm would be the fastest choice for X?

- A. Selection Sort
- B. Heapsort
- C. Mergesort
- D. Insertion Sort

## Observation: Insertion Sort on Almost Sorted Arrays

For arrays that are almost sorted, insertion sort does very little work.

- Left array: 5 inversions, so only 5 swaps.
- Right array: 3 inversion, so only 3 swaps.

A	E	E	L	M	O	B	R	X	Y	Z
A	E	E	L	M	O	B	R	X	Y	Z
A	E	E	L	M	O	B	R	X	Y	Z
A	E	E	L	M	O	B	R	X	Y	Z
A	E	E	L	M	O	B	R	X	Y	Z
A	E	E	L	M	O	B	R	X	Y	Z
A	E	E	L	M	O	B	R	X	Y	Z
A	E	E	L	M	O	B	R	X	Y	Z
A	E	E	L	M	O	B	R	X	Y	Z
A	E	E	L	M	O	B	R	X	Y	Z
A	E	E	L	M	O	B	R	X	Y	Z
A	E	E	L	M	O	B	R	X	Y	Z
A	E	E	L	M	O	B	R	X	Y	Z
A	E	E	L	M	O	B	R	X	Y	Z
A	E	E	L	M	O	B	R	X	Y	Z
A	E	E	L	M	O	B	R	X	Y	Z

A	E	L	E	M	O	P	R	X	Y	S
A	E	L	E	M	O	P	R	X	Y	S
A	E	L	E	M	O	P	R	X	Y	S
A	E	L	E	M	O	P	R	X	Y	S
A	E	L	E	M	O	P	R	X	Y	S
A	E	L	E	M	O	P	R	X	Y	S
A	E	L	E	M	O	P	R	X	Y	S
A	E	L	E	M	O	P	R	X	Y	S
A	E	L	E	M	O	P	R	X	Y	S
A	E	L	E	M	O	P	R	X	Y	S
A	E	L	E	M	O	P	R	X	Y	S
A	E	L	E	M	O	P	R	X	Y	S
A	E	L	E	M	O	P	R	X	Y	S
A	E	L	E	M	O	P	R	X	Y	S
A	E	L	E	M	O	P	R	X	Y	S
A	E	L	E	M	O	P	R	X	Y	S

# Picking the Best Sort (Poll Everywhere)

---

Suppose we do the following:

- Read 1,000,000 integers from a file into an array of length 1,000,000.
- Mergesort these integers.
- Select one integer randomly and change it.
- Sort using algorithm X of your choice.
- In the worst case, we have 999,999 inversions:  $\Theta(N)$  inversions.

Which sorting algorithm would be the fastest choice for X? Worst case run-times:

- A. Selection Sort:  $\Theta(N^2)$
- B. Heapsort:  $\Theta(N \log N)$
- C. Mergesort:  $\Theta(N \log N)$
- D. Insertion Sort:  $\Theta(N)$**

# Insertion Sort Sweet Spots

---

On arrays with a small number of inversions, insertion sort is extremely fast.

- One exchange per inversion (and number of comparisons is similar).  
Runtime is  $\Theta(N + K)$  where  $K$  is number of inversions.
- Define an ***almost sorted*** array as one in which number of inversions  $\leq cN$  for some  $c$ . Insertion sort is excellent on these arrays.

Less obvious: For small arrays ( $N < 15$  or so), insertion sort is fastest.

- More of an empirical fact than a theoretical one.
- Theoretical analysis beyond scope of the course.
- Rough idea: Divide and conquer algorithms like heapsort / mergesort spend too much time dividing, but insertion sort goes straight to the conquest.
- The Java implementation of Mergesort does this ([Link](#)).

# Sorts So Far

---

	Best Case Runtime	Worst Case Runtime	Space	Demo	Notes
<a href="#">Selection Sort</a>	$\Theta(N^2)$	$\Theta(N^2)$	$\Theta(1)$	<a href="#">Link</a>	
<a href="#">Heapsort</a> (in place)	$\Theta(N)^*$	$\Theta(N \log N)$	$\Theta(1)$	<a href="#">Link</a>	Bad cache (61C) performance.
<a href="#">Mergesort</a>	$\Theta(N \log N)$	$\Theta(N \log N)$	$\Theta(N)$	<a href="#">Link</a>	Fastest of these.
<a href="#">Insertion Sort</a> (in-place)	$\Theta(N)$	$\Theta(N^2)$	$\Theta(1)$	<a href="#">Link</a>	Best for small N or almost sorted.



# Shell's Sort (Extra)

## (Not on Exam)

## Optimizing Insertion Sort: Shell's Sort

---

Big idea: Fix multiple inversions at once.

- Instead of comparing adjacent items, compare items that are one stride length  $h$  apart.
- Start with large stride, and decrease towards 1.
- Example:  $h = 7, 3, 1$ .

h=1 is just  
insertion sort.

S O R T E X A M P L E

M O R T E X A S P L E	(1 swap, > 1 inversion)
M O R T E X A S P L E	(0 swaps)
M O L T E X A S P R E	(1 swap, > 1 inversion)
M O L E E X A S P R T	(1 swap, > 1 inversion)

M	O	L	E	E	X	A	S	P	R	T			
E	O	L	M	E	X	A	S	P	R	T	(1 swap,	>1 inversions)	
E	E	L	M	O	X	A	S	P	R	T	(1 swap,	>1 inversions)	
E	E	L	M	O	X	A	S	P	R	T	(0 swaps,	0 inversions)	
A	E	L	E	O	X	M	S	P	R	T	(2 swaps,	>1 inversions)	
A	E	L	E	O	X	M	S	P	R	T	(0 swaps,	0 inversions)	
A	E	L	E	O	P	M	S	X	R	T	(1 swaps,	>1 inversions)	
A	E	L	E	O	P	M	S	X	R	T	(0 swaps,	0 inversions)	
A	E	L	E	O	P	M	S	X	R	T	(0 swaps,	0 inversions)	

A	E	L	E	O	P	M	S	X	R	T
A	E	L	E	O	P	M	S	X	R	T
A	E	L	E	O	P	M	S	X	R	T
A	E	L	E	O	P	M	S	X	R	T
A	E	E	L	O	P	M	S	X	R	T
A	E	E	L	O	P	M	S	X	R	T
A	E	E	L	O	P	M	S	X	R	T
A	E	E	L	M	O	P	S	X	R	T
A	E	E	L	M	O	P	S	X	R	T
A	E	E	L	M	O	P	R	S	X	T
A	E	E	L	M	O	P	R	S	T	X

Total swaps: 14 (vs. 31)

# Shell's Sort: Generalization and Performance

$h=1$  is just normal insertion sort.

- By using large strides first, fixes most of the inversions.

We used 7, 3, 1. Can generalize to  $2^k - 1$  from some  $k$  down to 1.

- Requires  $\Theta(N^{1.5})$  time in the worst case (see CS170).
- Other stride patterns can be faster.

1, 3, 7, 15, 31, 63, ...	$\Theta(N^{3/2})$
1, 3, 5, 9, 17, 33, 65, ...	$\Theta(N^{3/2})$
1, 2, 3, 4, 6, 8, 9, 12, ...	$\Theta(N \log^2 N)$
1, 4, 13, 40, 121, ...	$\Theta(N^{3/2})$
1, 3, 7, 21, 48, 112, ...	$O(N^{1+\sqrt{8 \ln(5/2)/\ln N}})$
1, 8, 23, 77, 281, ...	$O(N^{4/3})$
1, 5, 19, 41, 109, ...	$O(N^{4/3})$

# Sorts So Far

	Best Case Runtime	Worst Case Runtime	Space	Demo	Notes
<a href="#">Selection Sort</a>	$\Theta(N^2)$	$\Theta(N^2)$	$\Theta(1)$	<a href="#">Link</a>	
<a href="#">Heapsort</a> (in place)	$\Theta(N)^*$	$\Theta(N \log N)$	$\Theta(1)$	<a href="#">Link</a>	Bad cache (61C) performance.
<a href="#">Mergesort</a>	$\Theta(N \log N)$	$\Theta(N \log N)$	$\Theta(N)$	<a href="#">Link</a>	Fastest of these.
<a href="#">Insertion Sort</a> (in-place)	$\Theta(N)$	$\Theta(N^2)$	$\Theta(1)$	<a href="#">Link</a>	Best for small N or almost sorted.
<a href="#">Shell's Sort</a>	$\Theta(N)$	$\Omega(N \log N)$ , $O(?)$	$\Theta(1)$	N/A	Rich theory!

An earlier version of this slide assumed that Heapsort was always given an array with no duplicate elements. If we omit this assumption, heapsort's best case is  $\Theta(N)$ .

# Citations

---

Lego man:

<http://thetechnicgear.com/wp-content/uploads/2014/02/sorting-lego.jpg>

Keynote Animations of Algorithms courtesy of Kevin Wayne